

PERFORMANCE-ORIENTED COMPUTING

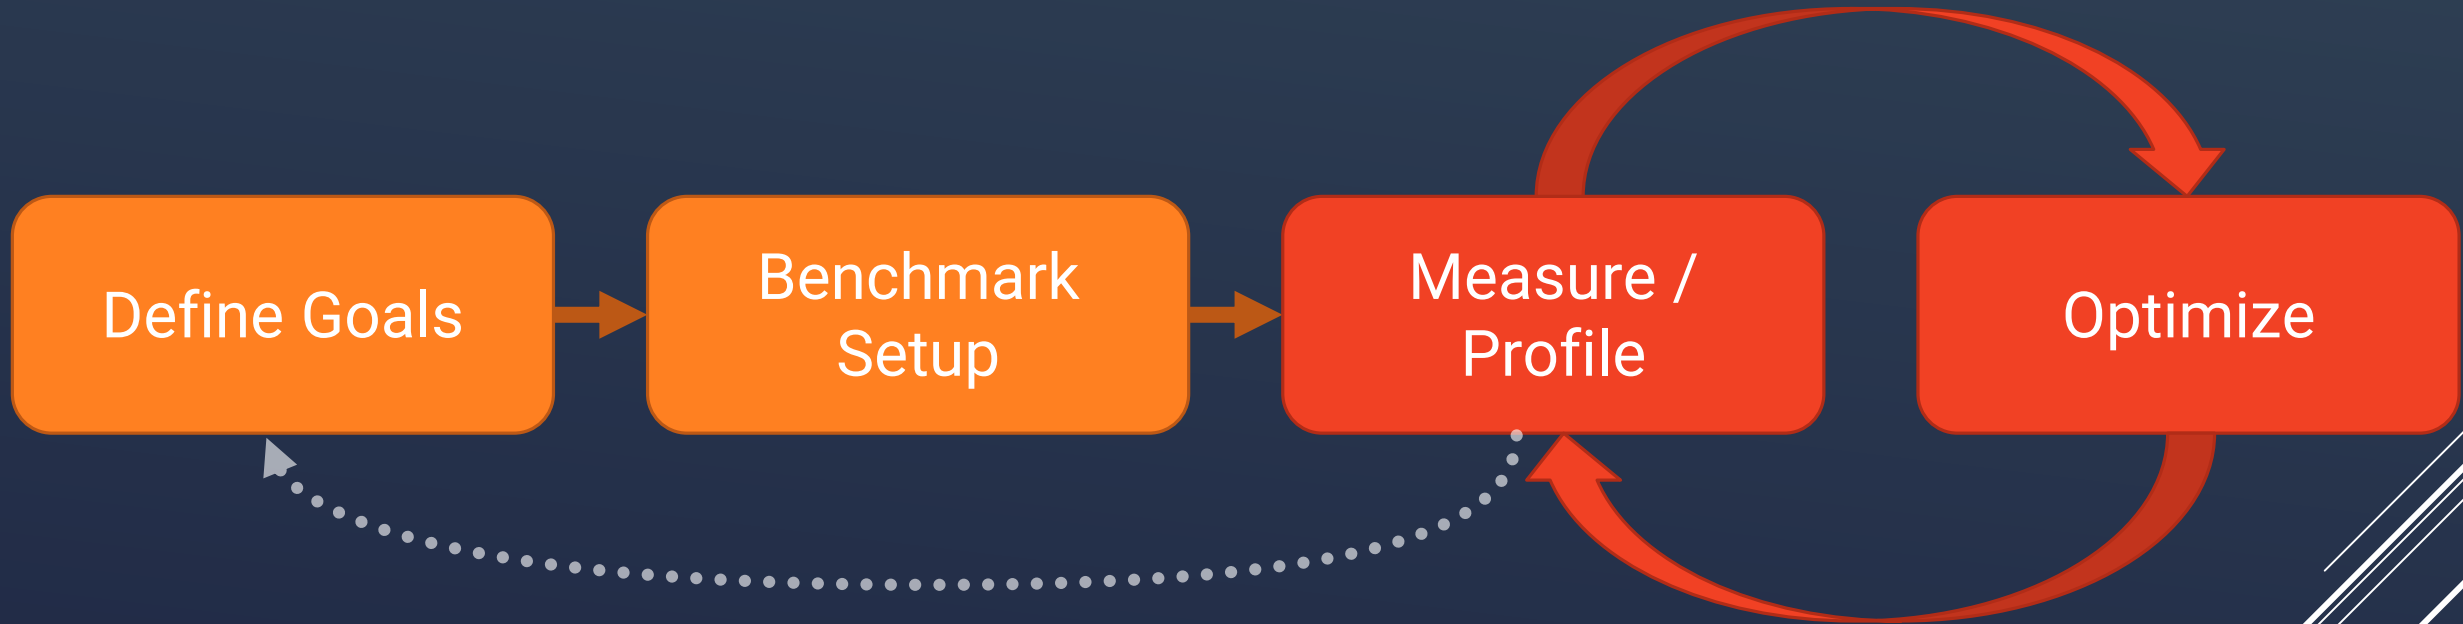
Optimization (0) - Overview



GOALS

- ▶ Before we get to the individual optimization topics, provide some general **definitions**
- ▶ Understand what “Optimization” **means**
- ▶ Be able to **classify optimizations** into 4 main categories

PERFORMANCE-ORIENTED DEVELOPMENT



- ▶ Optimization is a **cyclic** process, interleaving measuring / profiling with “actual” optimization work
- ▶ This aims to avoid *premature* or *ineffectual* optimization
 - ▶ We have the measurements to know **where** to effectively spend effort **before** we do so

DEFINING “OPTIMIZATION”

OPTIMIZATION – DEFINITION

- ▶ Exact definition varies with the Context of the optimization
- ▶ In the compiler context:

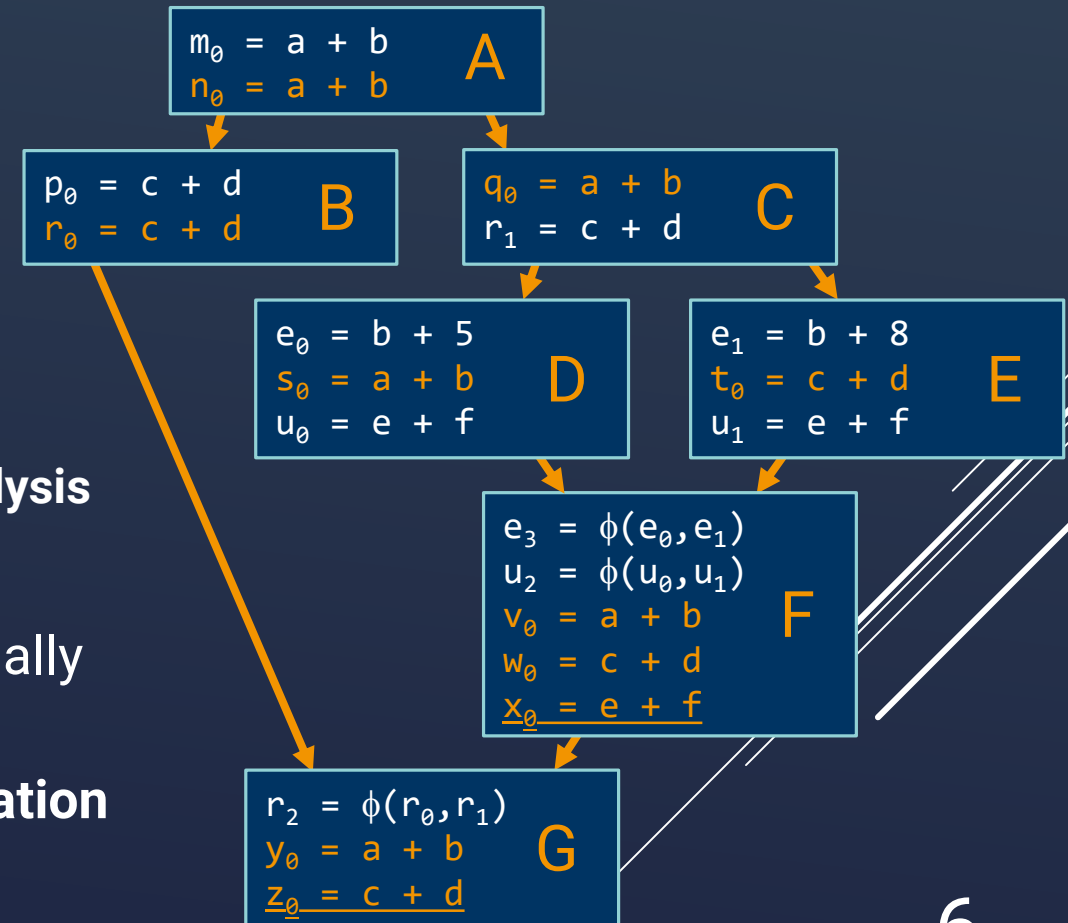
“Transformation of a program A into a program B, which has an identical input/output behavior but improves non-functional aspects.”

- ▶ Examples:
 - ▶ Redundant subexpression elimination
 - ▶ Loop-invariant code motion
 - ▶ Lots of other “classic” compiler optimizations
- ▶ This definition is very **safe**, but limits the space for optimizations

SIDENOTE: COMMON SUBEXPRESSION ELIMINATION


An expression **x (op) y** is redundant **if and only if**, along every path from the procedure's entry, it has been evaluated, and its constituent subexpressions (**x & y**) have **not** been redefined.

- ▶ Example: **colored** expressions are redundant
 - ▶ Some of these are relatively easy to find using **local** methods, while others require expensive **global analysis**
 - ▶ Difficulty of applying optimization affects its utility
 - ▶ Sometimes, performing this transformation is actually *slower*
- Both of these points **also apply to manual optimization** efforts to some extent!




LIMITS OF STRICT DEFINITIONS

- ▶ Compilers need a **strict** definition, they cannot judge if a change is “important”
- ▶ However, in practice, such a definition would **limit** our optimization space
- ▶ Changes in output may be irrelevant for **actual purpose**
 - ▶ *Example: Distance calculation in computer vision*



Result, FP32



Result, FP16

→ Different results (bitwise), but equivalent in practice!

ALTERNATIVE DEFINITION

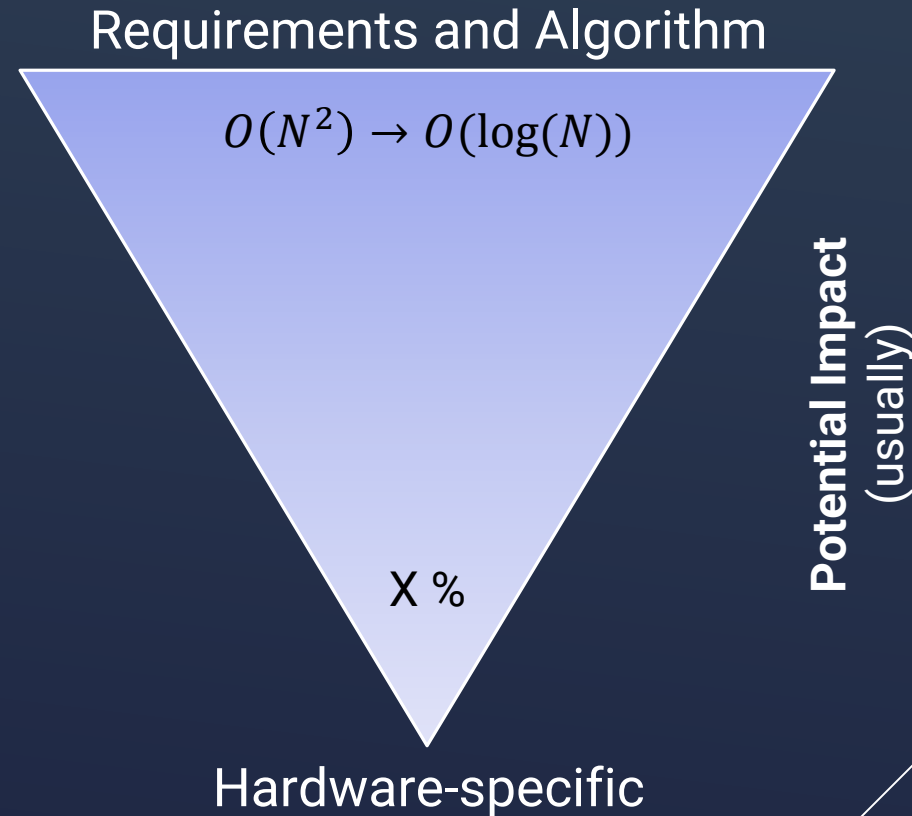
“Arbitrary changes in a program that improve non-functional aspects and still meet all other functional requirements.”

- ▶ I.e. exact equivalence on the individual bit level is **not** necessary
- ▶ Instead: Continued fulfillment of the specified **requirements**
 - ▶ Functional requirements on our program should be defined a priori (standard software engineering practices)
- ▶ *Example:* "Distance accuracy of the result is 1mm at a range of 10m".

CATEGORIES OF OPTIMIZATION

CATEGORIES OF OPTIMIZATION

1. **Adaptation of requirements**
2. **Algorithmic** optimization
3. **Procedural** optimization
4. **Hardware-specific** optimization



ADAPTATION OF REQUIREMENTS

- ▶ Improving performance or reducing resource consumption by **adapting requirements**
- ▶ The goal is to identify **non-essential requirements** that have a large impact on performance (compared to their necessity)
- ▶ Applicability is very application specific, and must be done in coordination with **domain experts**
- ▶ In contrast to other forms of optimization, **usually not objectively measurable and assessable**
 - ▶ Performance improvement can be measured, but not the influence on the overall quality – and thus the *meaningfulness* of the optimization

ADAPTATION OF REQUIREMENTS

- EXAMPLE -



- ▶ Alpha-blending has significant performance impact on Nintendo Switch (memory bandwidth), relatively minor visual impact
- Requirement adaptation: **reduction of background vegetation**

ALGORITHMIC OPTIMIZATION

- ▶ **Most important** category of optimizations

- ▶ By far the largest potential of possible performance improvements
- ▶ Wide range of subcategories

Potentially enables
reduction in complexity
Class!

(Usually not possible with
procedural or HW-specific
optimization)

- ▶ Includes all optimizations that achieve the same result in a *fundamentally different way*
- ▶ Algorithmic optimizations often also require **changes to data structures**

ALGORITHMIC OPTIMIZATION

- (SILLY) EXAMPLE -

Sorting

- ▶ If a program uses a Bubblesort or other sub-optimal sorting algorithm
- ▶ Use Quicksort or other efficient sort instead
- ▶ $O(N^2) \rightarrow O(N \log(N))$ complexity! (average case)
- ▶ This particular case is probably not a problem in most applications (nobody uses bubblesort)
- ▶ But general scenario extremely relevant!
 - ▶ Also in production software you frequently encounter ad-hoc solutions for problems for which there are **categorically better** algorithms

PROCEDURAL OPTIMIZATION

- ▶ Adaptations in the **implementation** of an **existing algorithm** to make the execution more efficient
- ▶ Distinction to Algorithmic Optimization not always 100% clear/binary
 - ▶ E.g. caching of existing results could be seen as procedural or algorithmic optimization
- ▶ Many of these are implemented as common compiler optimizations
 - ▶ Inlining, loop unrolling, ...
 - ▶ Common Subexpression Elimination (seen before) is another example

PROCEDURAL OPTIMIZATION

- EXAMPLE -

- ▶ Loop-invariant code motion

```
for(int i=0;...) {  
    ...  
    x[i] = y[i]  
        + (a*sin(b));  
    ...  
}
```



```
double c = a*sin(b);  
for(int i=0;...) {  
    ...  
    x[i] = y[i] + c;  
    ...  
}
```

- ▶ Implemented in all modern compilers in some form
- ▶ Like with many of these basic optimizations, challenge today is determining when to actually perform it
 - ▶ Alternative/opposite is called **rematerialization**

HARDWARE-SPECIFIC OPTIMIZATION

2 basic categories:

1. Changing program flow and/or data structures to work *better* on certain hardware
2. Leveraging special hardware units or features

► Important difference:

- **Category 1** *tunes* for certain target hardware, but remains generally executable
- **Category 2** *restricts* the portability of the program

Affects *performance*
portability

Affects plain old
portability

HARDWARE-SPECIFIC OPTIMIZATION

- EXAMPLES -

- ▶ For **category 1** – rather common: tuning for *cache hierarchies*
→ we'll discuss this in detail in the next chapter!
- ▶ For **category 2**: Counting the set bits in a given integer “n”

```
unsigned int count = 0;
while(n) {
    count += n & 1;
    n >>= 1;
}
```


```
unsigned int count =
    TBL[(n >> 0) & 0xff] +
    TBL[(n >> 8) & 0xff] +
    TBL[(n >> 16) & 0xff] +
    TBL[(n >> 24)]
```

```
unsigned int count =
    _mm_popcnt_u32(n);
```

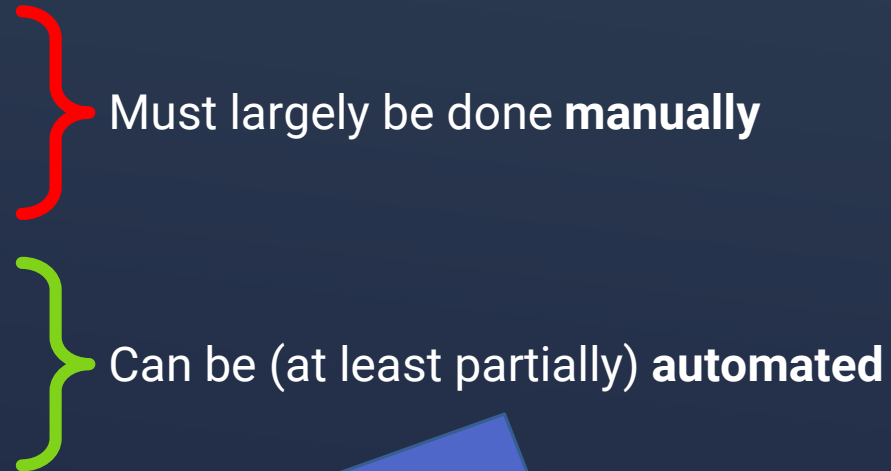
Performance		
Architecture	Latency	Throughput (CPI)
Icelake	3	1
Skylake	3	1
Broadwell	3	1
Haswell	3	1
Ivy Bridge	3	1

HARDWARE-SPECIFIC OPTIMIZATION

- ▶ Performance of hardware-optimized code is very **hardware-dependent**
- ▶ For explicitly hardware optimized code this seems obvious, but **also** optimizations that are *not obviously hardware specific* can show this behavior



CATEGORIES OF OPTIMIZATION AND TOOL SUPPORT

1. **Adaptation of requirements**
 2. **Algorithmic** optimization
 3. **Procedural** optimization
 4. **Hardware-specific** optimization
- 
- Must largely be done **manually**
- Can be (at least partially) **automated**

It makes sense to familiarize yourself with the relevant compiler features of your respective development platform.
(Not only standard optimizations, but also e.g. vectorization).

CONCLUSION

APPLICABILITY OF OPTIMIZATIONS

- ▶ **Adaptation of requirements** and **algorithmic** optimizations are **largely independent** of programming languages and software/HW platforms
 - ▶ But requirements adaptation is very domain-specific
 - ▶ **Procedural** and **hardware-focused** optimization differ depending on the programming environment and HW/SW stack
- We will primarily **focus on algorithmic optimization** in the remainder of this course!

SUMMARY

- ▶ The cycle of **performance-oriented development**
- ▶ **Definition** of “Optimization”
 - ▶ Strict compiler definition
 - ▶ More broadly applicable alternative & reasoning
- ▶ **Categories** of optimization
 - ▶ Adaptation of requirements
 - ▶ **Algorithmic optimization**
 - ▶ Procedural optimization
 - ▶ Hardware-specific optimization

QUESTIONS ?

