

PERFORMANCE-ORIENTED COMPUTING

Optimization (2) – Data Structures



GOALS

- ▶ Understand which structures are favorable for which **types of operations**
 - ▶ Know which **decision criteria** there are, and how to apply them
- ▶ Distinguish between the **theoretical complexity** of data structure operations, and their **real-world performance**
- ▶ Gain an **overview** of **general-purpose data structures** from a performance perspective
- ▶ An outlook on **specialized** data structures

OPERATIONS & DECISION CRITERIA

3

INTRODUCTION

- ▶ Selecting the proper data structure(s) for a specific purpose is an **essential part of algorithmic optimization**
- ▶ Selecting a suitable data structure may:
 1. Allow performing common (in your application) operations at a **reduced complexity class** compared to a less suitable structure
 2. Improve the **memory layout** and/or allow the memory hierarchy of your target platform to work more effectively, significantly improving performance

HOW TO DETERMINE WHETHER A DATA STRUCTURE IS SUITABLE?

Several common decision criteria:

- ▶ **Type** and **quantity** of data which needs to be stored
- ▶ **Access patterns**
 - ▶ **Where** do we perform **which types** of operations?
- ▶ Target **hardware** properties

Generally **not as critical as the first two points**, but can be a **tie-breaker**.

TYPES OF DATA

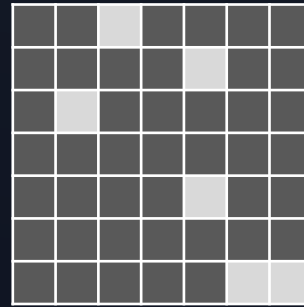
The data type can influence / limit the viable data structures:

- ▶ Are data elements **comparable**? Is there an ordering?
 - ▶ Required for e.g. many tree-based data structures

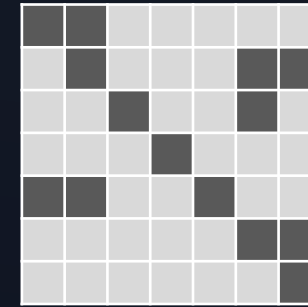
Might be used internally to implement data structures which are not semantically trees (e.g. sets, maps)!

- ▶ Is a **hash** operation defined on the data elements?

- ▶ Are elements **countable**?
 - ▶ If so, is there a **dense embedding** in some iteration space?



Dense Embedding



Sparse(-ish) Embedding

QUANTITY OF DATA

Both **total quantity** and the **size of individual elements** influence our choice of data structures

- ▶ **Total quantity:**

- ▶ **Low** total size → we might want to “waste” some space to save time
- ▶ **High** number of elements → **complexity class** of operations becomes more important than constant overheads or HW suitability

- ▶ **Element size:**

- ▶ **Small** individual elements need to be stored **densely** in memory
- We'll talk about a concrete example of this later in this lecture

ACCESS PATTERNS

We generally **don't care** about patterns happening uncommonly, e.g. only during initialization of a long-running program.

- ▶ Adapting our data structure selection to the *common access patterns* of our application can have significant performance advantages

Categorize by:

- ▶ **Type** of access:
 - ▶ Read, write, insert, delete, replace, search, ...
- ▶ **Position** of access:
 - ▶ Front, back, sequential, arbitrary, ...
- ▶ **Parallel** access:
 - ▶ From how many threads? Multiple readers/writers/both?

ACCESS PATTERN EXAMPLE

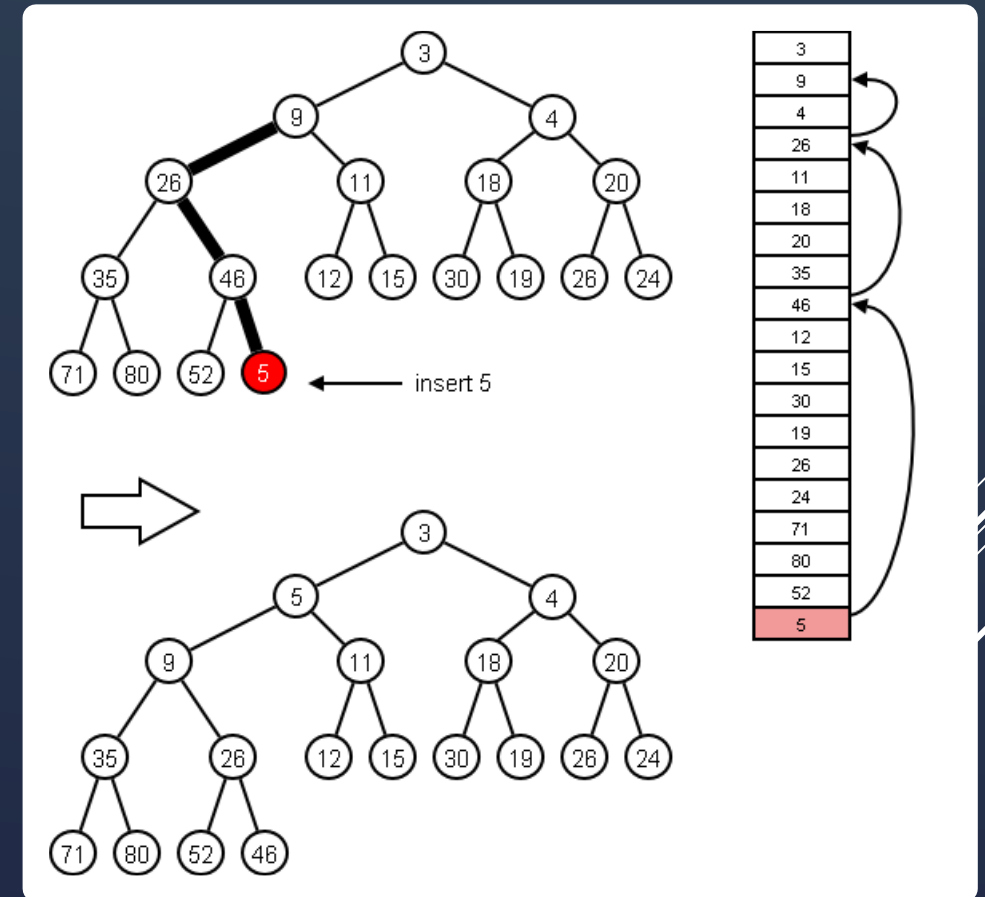
We want to arbitrarily add elements, but always access the one with the highest *priority*

- ▶ First idea: need data structure which we can **efficiently sort**

- ▶ *Even better*: one which is **always sorted**

→ **Priority Queue**

→ Still open implementation choices, e.g. How to store elements



TARGET HARDWARE CONSIDERATIONS

- ▶ **Additional factor** in the selection of data structures *after* algorithm considerations
- ▶ Importance depends on the **flexibility** of the target HW
 - ▶ i.e. we need more consideration towards HW when selecting a data structure for use on a GPU than a general-purpose CPU
- ▶ Most obvious constraint: **total memory**
 - ▶ Influences how much we need to focus on memory space efficiency

TARGET HARDWARE CONT.

Other relevant hardware aspects:

- ▶ **Caching / memory hierarchy**

- ▶ Depending on the sophistication and importance of the memory hierarchy, we might prefer data structures with fewer indirect accesses

- ▶ **Impact of branching**

- ▶ Much more impactful on e.g. GPUs than high performance CPUs

May sound obvious / straightforward, but generally we don't make data structures more indirect or branchy just because it's fun.

Often a **tradeoff between operational/algorithmic efficiency and hardware efficiency!**

PERFORMANCE THEORY VS PRACTICE IN DATA STRUCTURES



12

PERFORMANCE THEORY

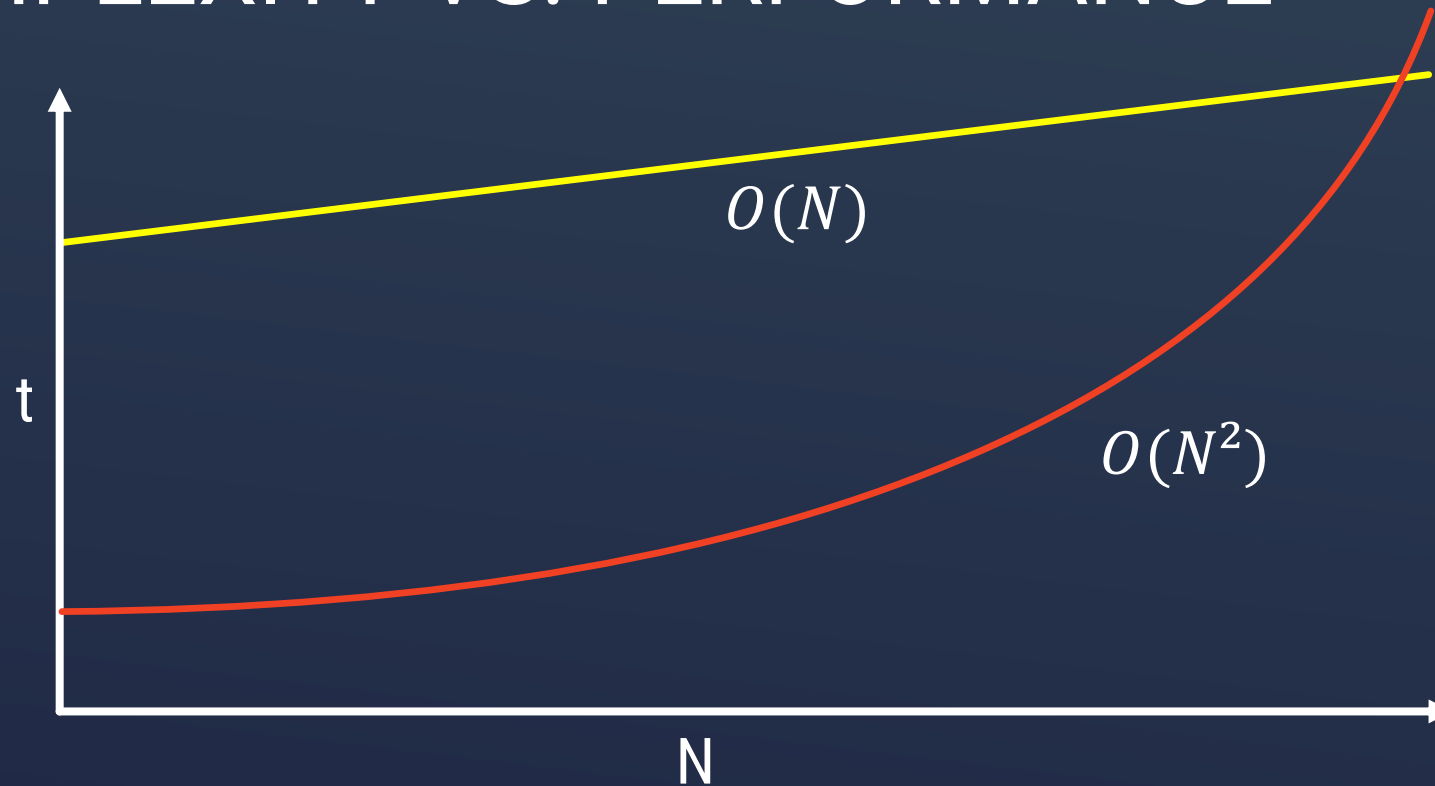
In **theory**, we generally express the expected performance of data structures by specifying the asymptotic bound of each potential operation using **big-O notation**

- ▶ Crucially, when we specify a complexity class, we **ignore constant factors**
 - ▶ This is often sensible and appropriate; but sometimes it isn't
- ▶ Also, it's important to distinguish between **average-case** and **worst-case** analyses

PERFORMANCE PRACTICE

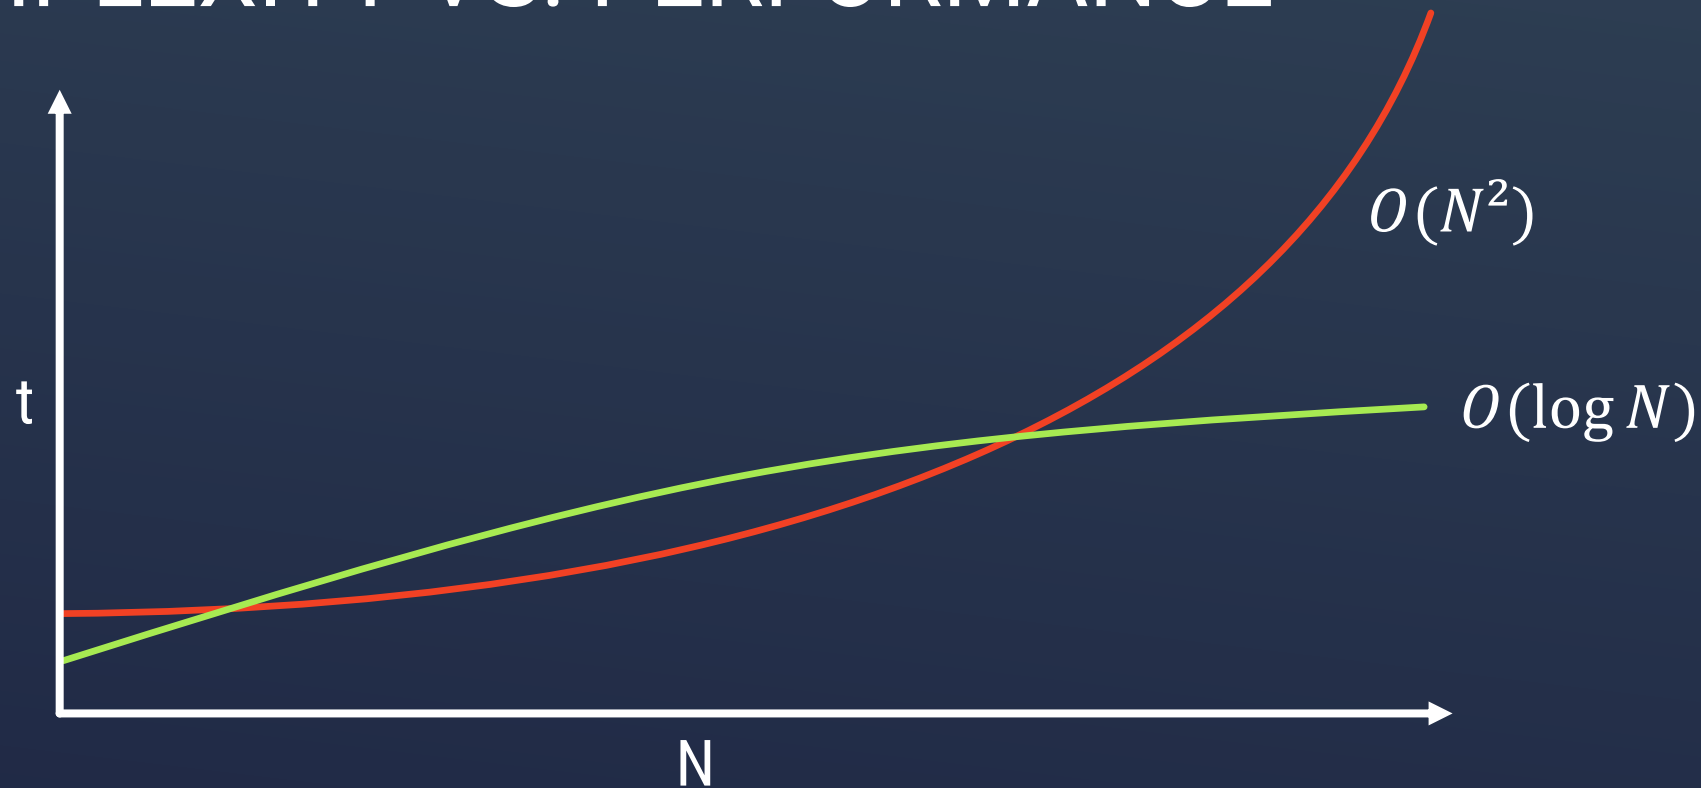
- ▶ In practice, **constant factors** – and even constant overheads – **can** make a large difference
 - ▶ And they are frequently **influenced by hardware factors**
- ▶ Rather obvious when the operations have the **same** complexity class
 - ▶ But depending on the degree of overhead and the expected N , *even differences in asymptotic complexity might not matter*
 - ▶ **Important:** decisions made based on this need to take into account the **expected evolution of N** , and might need to be revisited!

COMPLEXITY VS. PERFORMANCE



- *Need to be aware of:*
how much overhead (if any) there is → **where** the transition happens

COMPLEXITY VS. PERFORMANCE



- *Need to be aware of:*
where our program will be on the **N** scale

OPERATION MIX

- ▶ In real programs, your interaction with a given data structure is almost **never a single operation**
- ▶ Have to take into account **aggregate** performance of ***operation mix***
- ▶ Also important to consider distinct **phases** of a program
 - ▶ I.e. expensive building of a data structure justified if it happens **once** at the start and is then queried very **frequently**

THE DANGERS OF INTUITION

- Between **array-like** lists and **linked lists**, which data structure would you expect to perform best in the following cases?

Insertions/Deletions	Reads	Elem. Size	N
1%	99%	8 Byte	100
10%	90%	8 Byte	100
50%	50%	8 Byte	100
1%	99%	512 Byte	1000
10%	90%	512 Byte	1000
...			

➔ We'll explore the answer to this in the exercises!

OVERVIEW OF SOME DATA STRUCTURES

And their performance properties

ARRAY-LIKE LISTS



- ▶ Perhaps the most basic data structure
- ▶ **Great for**
 - ▶ **Iterative traversal**
 - ▶ **Random indexed access**
 - ▶ Insertion at the end ...
 - ▶ ... as long as there is some extra space
- ▶ **Dense storage**, no extra metadata beyond size
 - ➔ Very space-efficient, **great cache properties**

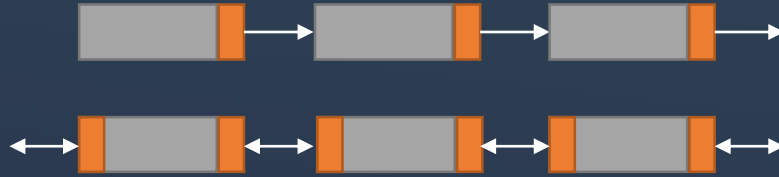
▶ **Weaknesses**

- ▶ Frequent arbitrary **insertion/deletion** (not at end)
- ▶ **Search**
 - ▶ Can be solved by sorting (if only searching by one criterion)
- ▶ **Pointer invalidation** / lack of reference stability

Pointers/references to individual elements might be **invalid after operations** on the data structure (as elements were reallocated).

Also known as **iterator invalidation** e.g. in C++.

LINKED LISTS



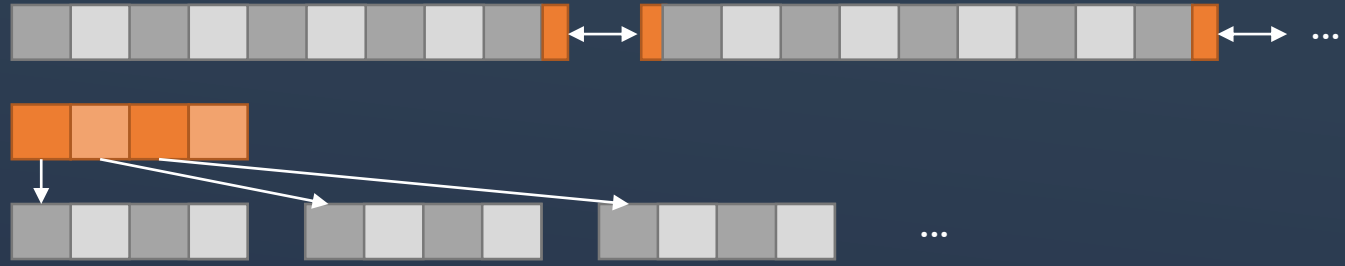
- ▶ Either single- or double link, depending on use case
- ▶ **Great for**
 - ▶ **Large** amounts of **data** (esp. large individual elements)
 - ▶ **Frequent** arbitrary **insertion/deletion**
- ▶ Provides **stable references**/pointers to elements

▶ **Weaknesses**

- ▶ **Not dense** in memory
 - slower traversal (and operations in general) due to caching and indirection
- ▶ No random indexed access
- ▶ Much less predictable for the compiler
 - lower optimization potential

The performance impact of these weaknesses is easy to underestimate!

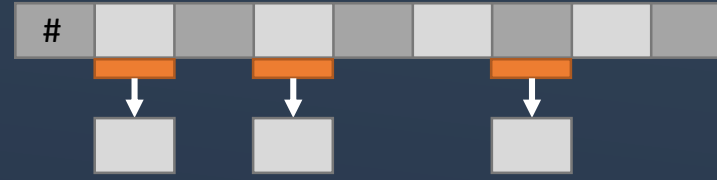
HYBRID LISTS



- ▶ Many data structures that seek to **combine** (some of!) the benefits of array-like and linked lists
- ▶ Examples:
 - ▶ Linked lists of fixed-size multi-element chunks
 - ▶ Array of pointers to fixed-size chunks
 - ▶ ...
- ▶ Can provide e.g. faster arbitrary insertion/deletion than arrays with better cache behaviour than linked lists
 - ▶ But always a **tradeoff**

A real-world example of this is the C++ `std::deque` data structure.

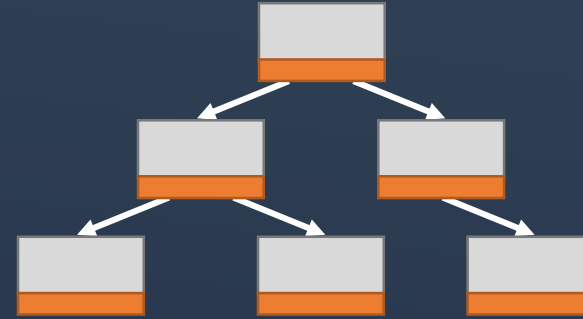
HASHSETS/MAPS



- ▶ Ideal for very **fast lookup** of specific elements
 - ▶ $O(1)$ when everything goes well
- ▶ **Requirement:** elements must be **efficiently hash-able**
- ▶ Also helpful, but not necessary: **prior knowledge** of roughly **how many** elements will be stored
 - ▶ Alleviates the need to change the number of buckets or re-hash

- ▶ Several performance-relevant details to hash out:
 - ▶ Open or closed addressing
 - ▶ Initial size
 - ▶ Tradeoff between collisions and hash function
- ➔ No change to fundamental properties / use cases

TREE STRUCTURES

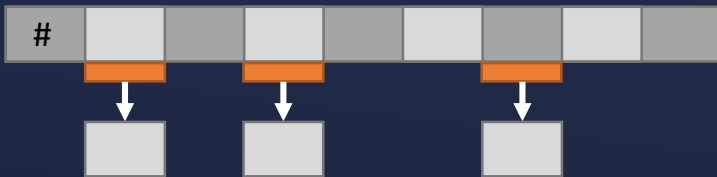


- ▶ Flexible for **many kinds of lookup / search purposes**
- ▶ Great when **sorted traversal** is frequently performed
- ▶ **Requires an ordering** on items
- ▶ **Disadvantage:** generally requires **balancing** operations in order to maintain performance after many insertions / deletions

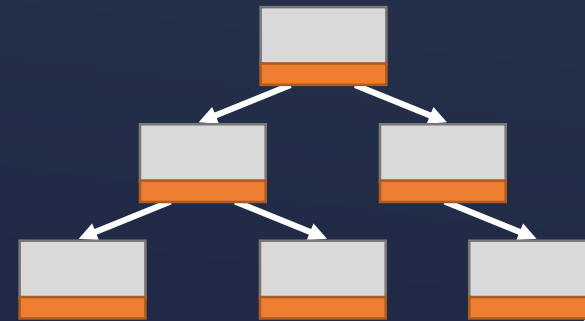
TREES VS HASHES

Both data structures are good for fast data retrieval. How to decide? – Comparison:

- ▶ Very fast **lookup**
- ▶ Potential for **performance degradation**
Usually $O(1)$, but can degrade to $O(n)$
- ▶ Comparatively easier to implement correctly
- ▶ Tradeoff between memory efficiency and performance



- ▶ Very fast sorted **traversal** (since already sorted)
- ▶ **Predictable** performance
 $O(\log(n))$ for most operations
- ▶ Can be complex to implement (esp. balancing)
- ▶ Predictable memory utilization



RING BUFFERS



- ▶ Very **efficient insertion and removal** of elements at **start/end**
- ▶ With **pointer stability!**
- ▶ Fast **traversal**, indexed lookup
- ▶ Very well suited for transporting data in a **parallel producer/consumer** scenario
 - ▶ Only need to update 2 indexes/pointers, can be performed atomically

- ▶ ***Disadvantages / constraints:***

- ▶ Maximum **number of entries** needs to be known **a priori**
- ▶ **Fixed memory allocation** of all elements, regardless of actual use

Relatively small *number* of constraints, but they are very significant for many use cases.

FLYWEIGHT PATTERN



- ▶ Useful when there are a relatively **small number of distinct**, complex **elements**
 - ▶ But they are arranged in **complex relationships** with many individual instances (e.g. graphs)
 - ▶ Represent them by flyweight instances which refer to the actual elements
- ▶ Similar to using pointers, but can deal with instances as **values**
 - ▶ Less error prone memory management
 - ▶ Potentially better performance

BEYOND THIS SELECTION

- ▶ There are many more *relatively* general purpose data structures
 - ▶ Bitsets
 - ▶ Interval containers
 - ▶ Bimaps
 - ▶ ...

➔ When there is a performance issue in your application that relates to operations on a container / data structure, determine the **decision criteria** we discussed (access patterns etc.) and then **search for a good fit!**

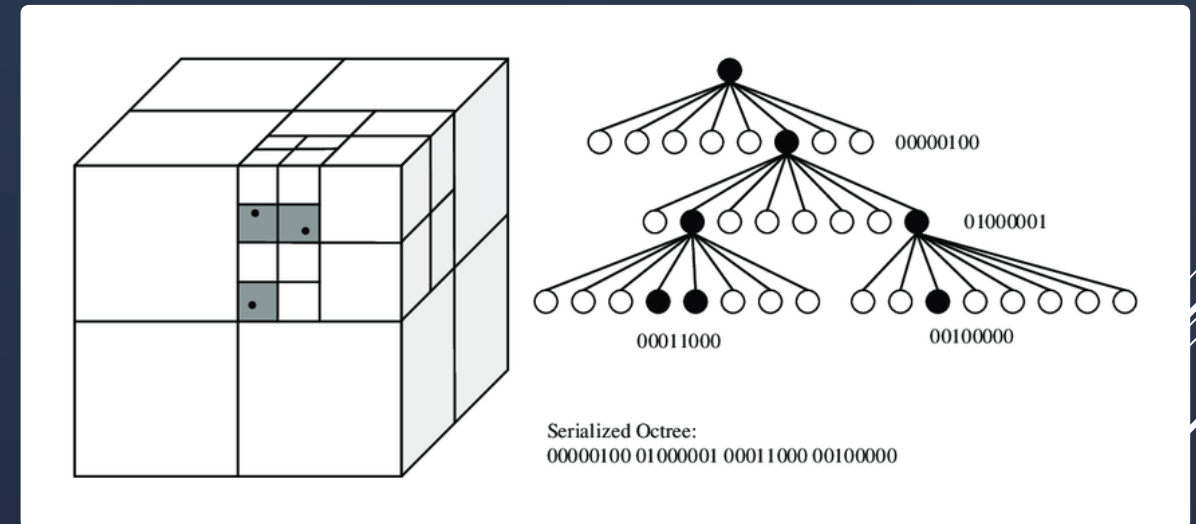
SPECIALIZED DATA STRUCTURES

WHAT DOES “SPECIALIZED” MEAN?

- ▶ Here, “specialized” means that these structures are built to either store one **particular type of data**, or support **particular kinds of operations**
- ▶ Especially in the context of a specific **application domain**
- ▶ Like many of these distinction, there is **no hard line** between “general” and “specialized”
 - ▶ E.g. interval containers or the flyweight pattern mentioned previously are clearly *more* specialized than a basic list, but still *less* specialized than some of the examples we’ll look at now

SPACE PARTITIONING STRUCTURES

- ▶ One example of a class of domain-specific data structures
 - ▶ Specialized for **queries in 3D space**
- ▶ Includes
 - ▶ **Octrees**
 - ▶ Kd-trees
 - ▶ Portal graphs
 - ▶ ...



➔ When a domain is widely explored in a performance-focused context then there is often a **wealth of data structures** for the same (or very similar) purpose!

HARDWARE SUPPORT FOR SPECIFIC DATA FORMATS

- ▶ Sometimes, domain-specific data formats are sufficiently **critical in a domain** to reach widespread **hardware support**
- ▶ In such cases, it's almost always advantageous to make use of these formats
- ▶ Hardware implementation is generally:
 - ▶ **Faster**
 - ▶ More **efficient** (in terms of energy use)
 - ▶ Very likely to be **correct**

HARDWARE SUPPORT EXAMPLE

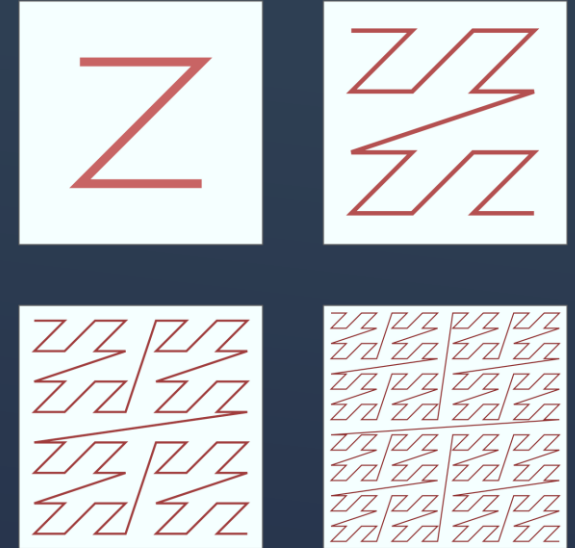
VIDEO AND AUDIO



- ▶ **Media files** are perhaps the most well-known data formats with widespread hardware acceleration
- ▶ E.g. all kinds of devices, from PCs over mobile phones to TVs support a variety of **audio and video codecs** in hardware
 - ▶ For several of them, allowing playback of files that wouldn't even be possible in software on the same platform
- ▶ In terms of performance, it might well be worth it to e.g. use a “worse” format with HW support rather than a “better” format in software
- ▶ Potential issues with **licensing**, but that's not a performance constraint
 - ▶ Unencumbered high-end formats now starting to become available in HW!

HARDWARE SUPPORT EXAMPLE TEXTURES & COMPRESSION

- ▶ Textures on GPUs stored with **2D locality**
 - ▶ Uses a mapping to memory which follows a **space-filling curve**
 - ▶ Usually **Z-curve** due to efficient implementation
 - ▶ Highly **efficient** access, texture caching and filtering in HW
- ▶ Also HW support for various **block compression formats**
 - ▶ BC1 – BC7 (as named in D3D)
 - ▶ Also some less commonly supported formats like
ASTC (Adaptive Scalable Texture Compression – primarily mobile)



CONCLUSION

35

SUMMARY

► Operations & Decision Criteria

- Types of Data
- Quantity of Data
- Access Patterns
- Hardware Considerations

► Overview of Data Structures

- Array-like Lists
- Hashsets/maps
- Linked Lists
- Tree structures
- Hybrid Lists
- Trees vs. Hashes

► Performance Theory vs Practice

- The impact of **constant factors** and overhead *before* the asymptotic limit
- Operation Mix
- The Dangers of Intuition

- Ring Buffers
- Flyweight Pattern
- Specialized Data Structures

QUESTIONS ?

