

# PERFORMANCE-ORIENTED COMPUTING

Optimization (3) – Algorithms



# GOALS

- ▶ Learn general **strategies** for algorithmic optimization
  - ▶ Optimizing individual algorithms isn't all that meaningful – the challenge is running into a **problem in some custom algorithm in software** you need to optimize
  - ▶ Try to provide generally applicable strategies to deal with this issue
- ▶ Discuss the potential of **memoization** for many algorithms
  - ▶ As an example of the general tradeoff between space and time

# INTRODUCTION & OVERVIEW

3

# OVERVIEW

- ▶ As we discussed before, algorithmic optimization is the **most important** category of optimization
  - ▶ Largest potential of performance improvement (without adapting requirements)
  - ▶ Broad spectrum of sub-categories
- ▶ Contains all optimizations which solve a problem in a **fundamentally different** way

*Most optimizations which touch data structures that we discussed before **are** algorithmic optimizations!*

# CHALLENGES

- ▶ In practice, algorithmic optimization is **very challenging**
  - ▶ **Widely-used** and easily identified algorithms are **already optimized**
    - May need to tune and select for specific target hardware
- ▶ *What if the algorithm **isn't** widely used or easily identified?*
  - ▶ Need to learn to recognize algorithmic similarities
  - ▶ Sometimes, *might* actually need to come up with a new algorithm
    - ▶ But more likely, a heuristic **based on** an existing algorithm and tuned for your use case

# STRATEGIES FOR ALGORITHMIC OPTIMIZATION

1. Select **appropriate data structures** for your operations  
→ We already discussed this selection process in depth in the previous chapter
2. Apply existing **best practices**
3. **Leverage** program-specific information and **context**



Some hints for these  
provided in this lecture.

# ALGORITHMIC OPTIMIZATION STRATEGIES

And Some Examples



7

# APPLYING BEST PRACTICES

► Essentially three steps:

1. **Recognize** program task/process as *instance of an existing problem*
2. **Study literature** to find best approach
3. **Implement** – or better, find a suitable library or existing implementation

This is often the hardest part!

Can enable you to **benefit from** years or even decades of algorithm **research, optimization**, and in the case of libraries even implementation experience and **polish**.



# HOW TO RECOGNIZE INSTANCES OF WELL-STUDIED ALGORITHMS

- ▶ Familiarize yourself with important algorithms, especially in **optimization** and **graph theory**
  - ▶ E.g. knapsack problem, maximum flow, ...
- ▶ **Analyze the problem domain** and the nature of the data involved
  - ▶ Look for patterns, similarities, or **common structures** with known algorithms
  - ▶ Compare problem **characteristics**: objectives, constraints, input/output, ...

This is a skill that improves with **experience** and exposure to various problem domains. Practice and continuous learning will enhance your ability to identify and apply appropriate techniques.

# EXAMPLE: REGISTER ALLOCATION IN COMPILERS


- ▶ Essential part of the **code generation** phase of compilers
- ▶ **Goal:** use the limited set of available ISA registers as efficiently as possible to store the (potentially much larger) set of intermediate values
- Can be mapped to the **graph coloring** problem (specifically *vertex coloring*)
  - ▶ Limited set of resources (colors) need to be assigned to graph vertices
  - ▶ No two directly connected vertices may share the same color



# EXAMPLE:

## REGISTER ALLOCATION IN COMPILERS


- ▶ **Goal:** use the limited set of available ISA registers as efficiently as possible to store the (potentially much larger) set of intermediate values
- ▶ Make use of an *interference graph*
  - ▶ **Vertices:** temporary values to be stored
  - ▶ **Edges:** connect values required concurrently



# EXAMPLE:

## REGISTER ALLOCATION IN COMPILERS


- ▶ **Goal:** use the limited set of available ISA registers as efficiently as possible to store the (potentially much larger) set of intermediate values
- ▶ Make use of an *interference graph*
  - ▶ **Vertices:** temporary values to be stored
  - ▶ **Edges:** connect values required concurrently
- ▶ If graph can be colored with  $N$  colors  
→ values can be kept in  $N$  registers!
- ▶ If fewer registers available  
→ need *spilling* decision  
→ **domain-specific heuristics**



# EXAMPLE: REGISTER ALLOCATION IN COMPILERS

- ▶ *How would you think of this approach?*
- ▶ At least three important components:
  1. **Know about the graph coloring problem**  
→ *"Familiarize yourself with important algorithms"*
  2. **Understand the core constraints at an abstract level**  
(i.e. value lifetimes and their interaction with limited register slots)  
→ *"Analyze the problem domain"*
  3. **Formalize the problem in a structured way**  
(i.e. in this case as a graph)

→ **General strategy** which applies to many problems and domains



# LEVERAGING PROGRAM-SPECIFIC CONTEXT

- ▶ The optimal **generic** algorithm may not be ideal in a **specific** situation / context in a given optimization target program
- ▶ Knowledge of the relevant **context** can be leveraged
  - ▶ E.g. range of input data, structure of objects, semantic constraints on results, ...
- ➔ In practice, this information is often explicitly or implicitly used in application- or **domain-specific heuristics**

# EXAMPLE: A\* PATHFINDING

- ▶ Dijkstra's Algorithm is **optimal** for the case of **generic** graphs
- ▶ However, we can achieve far better average-case performance in common **special cases** by **leveraging additional information!**
- ▶ **A\*** uses the additional concept of a minimum possible distance from a given node to steer the search for the shortest path



# EXAMPLE: A\* PATHFINDING


- ▶ For this example:
  - ▶ Assume that graph is embedded in a conventional 2D space exactly matching the layout in the illustration
  - Use **Euclidian distance** as minimum!
- ▶ **Important:** A\* is not a “better algorithm” in a general sense
  - **Context** is essential!





# EXAMPLE: A\* PATHFINDING

- ▶ Heuristics are only as good as their assumptions
  - ▶ There are usually degenerate cases which fall back to worst-case complexity of the underlying algorithm
  - ▶ On the right you have an example of that for A\*
- ▶ **Again:** understanding the exact **context** and statistical distribution of various scenarios in **your application** is essential



# A NOTE REGARDING THE EXAMPLES

- ▶ You might have noticed that both of the examples here are based on algorithms in graph theory
- ▶ This is **not** because everything you'll encounter will reduce to some algorithm from graph theory with an efficient solution
  - ▶ It's primarily for educational purposes, because these graph algorithms are relatively easy to effectively **visualize** and explain

*That said, a **lot** of problems **do** map really nicely to graph algorithms, and it is absolutely worth it to study the most important graph algorithms!*

# SUMMARY ALGORITHMIC OPTIMIZATION

- ▶ **Mapping** to well-studied problems
  - ▶ Perhaps the most **challenging** part, especially when it comes to less widely used algorithms
  - ▶ Requires knowledge of a **broad set** of common algorithms and approaches
  - ▶ Improves incrementally with **experience**
- ▶ Leveraging **program-specific context**
  - ▶ Generally **requires** knowledge of the domain
  - ▶ Can make sense to discuss this with domain specialists

→ In both cases, a **holistic view** of the program and its context is necessary!

# MEMOIZATION

And the tradeoff between space and time



20

# OVERVIEW

- ▶ **Memoization** is an optimization technique which can be extremely effective
- ▶ Basic idea:

**Store** the results of expensive function calls and **reuse** them when the **same inputs** are encountered again.

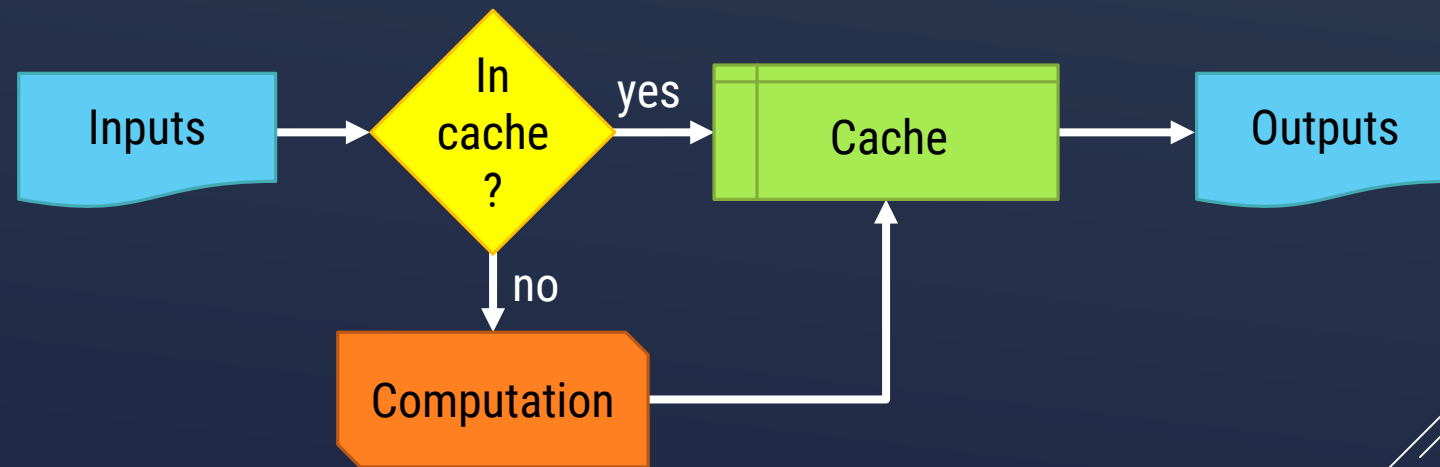
- ▶ Avoids **redundant** computations **at runtime**  
(rather than e.g. redundant expression elimination in the compiler)
  - ▶ Unlike redundancy elimination in the compiler, memoization can be applied with dynamic inputs and much more complex control flow

# OVERVIEW

## No Memoization



## With Memoization




# KEY CONSIDERATIONS

- ▶ **Most important:** memoization requires *pure functions*
  - ▶ I.e. no side effects, no access to any inputs other than arguments
  - Same inputs must always produce the same outputs
- ▶ Especially for larger outputs, **cache invalidation** / **cleaning** strategies need to be employed
  - ▶ This can be very complex and increase the implementation effort while reducing effectiveness

# CACHING STRATEGIES

- ▶ Primarily requires a **fast lookup** structure / operation
- ▶ Generally realized using an **associative array/map** backed by a **hash table**
  - ▶ Special cases might use a more algorithm-specific data structure, e.g. a simple dense array or specialized tree-based map
  - ▶ Input parameters serve as keys, output stored as values





# PARALLELIZATION

- ▶ **Note:** pure function calls can generally be executed freely and trivially in parallel **without** synchronization
- ➔ Caching can break this property!
- ▶ *Possibilities:*
  - ▶ **Per-thread Cache**
    - ▶ Requires even more memory, but retains synchronization-free property
  - ▶ **Read/write locked shared data structure**
    - ▶ Some overhead, but threads can share progress / results

**Best choice** depends on speed/**frequency** of individual computations, **size** of the input/output cache, and the expected **degree of parallelism**.

# IMPLEMENTATION OPTIONS

*Major options:*


## 1. **Manually add caching** logic within the function

- ▶ *Advantages:* potentially highest performance
- ▶ *Disadvantages:* mixes actual algorithm logic with memoization implementation; harder to maintain and tune individually

## 2. **Wrap the function** in an interface which adds caching

- ▶ *Advantage:* Clean separation; easy to read, maintain and tune
- ▶ *Disadvantage:* might introduce mutual dependence in recursive case

In **both cases** a memoization **library or framework** can be used. Some languages might support simple function decorations to enable memoization.




# SPECIAL CASE: RECURSIVE DIVIDE-AND-CONQUER

- In **recursive divide-and-conquer** algorithms, memoization can be an **essential** part of the overall algorithm!



No memoization



With memoization

# SPECIAL CASE: RECURSIVE DIVIDE-AND-CONQUER

- ▶ In **recursive divide-and-conquer** algorithms, memoization can be an **essential** part of the overall algorithm!
- ▶ Depending on the structure of data and degree of reuse, can **substantially reduce time complexity**
  - ▶ E.g. from exponential to quadratic or even linear in specific cases!
- ▶ This is one reason why **functional languages** sometimes have more first-class memoization support
- ▶ Particularly important application: **Dynamic Programming**
  - ▶ In this context, sometimes also called **tabulation**, mostly depending on how results are stored

# MEMOIZATION TRADE-OFFS

- ▶ *Already discussed*: need for pure functions, **no side effects**, potential impact on ease and **effectiveness of parallelization**

Other performance / suitability considerations:

- ▶ **Size/variety** of the set of expected **inputs**
  - ▶ Larger input set  
→ lower cache hit rate and more overhead for cache structure
- ▶ **Space vs. Time**
  - ▶ Larger keys and/or values  
→ needs longer computations in order for caching to pay off  
→ obviously requires more memory, and potentially more memory management

# CONCLUSION

30

# SUMMARY

## ▶ **Algorithmic Optimization**

- ▶ Huge potential for performance improvement
- ▶ Often challenging to apply

## ▶ **Strategies**

- ▶ Applying best practices
- ▶ Leveraging context and domain knowledge

→ Study **common algorithms** and gain a **holistic overview** of the program you are working on and its domain!

## ▶ **Memoization**

- ▶ Idea & key considerations
  - ▶ Pure functions
- ▶ Caching strategies
- ▶ Parallelization
- ▶ Implementation options
- ▶ Trade-offs

## ▶ **Special Case: Recursion**

- ▶ Dynamic Programming

QUESTIONS ?

